

Lecture #23

CS 170

Spring 2021



Streaming Algorithms

- Motivation

- Vast amounts of data "streaming" by, too much to store
 - Search engine tracking clicks on websites
 - Router monitoring network traffic
 - Data arriving from sensors
- Is there a much (exponentially) smaller data structure that we can quickly update on the fly, and query when needed?
- Monte Carlo algorithms: fast, probably accurate

3 Examples of Streaming Algorithms

- Simple: Counting Total Sales
 - Input: n sales with prices p_1, p_2, \dots, p_n
 - Desired output: $P = \sum_{i=1}^n p_i$
 - Initialize $C=0$, Update $C=C+p_i$, Query: return C
 - Memory Requirement: $\lceil \log_2 P \rceil$
- Morris's Alg. for Approximate Counting
 - Like above, with $p_i=1$
 - Goal: use $O(\log_2 \log_2 n)$ bits, not $\lceil \log_2 n \rceil$
- Flajolet & Martin (FM) Alg. for Distinct Elements
 - Count # distinct integers among i_1, \dots, i_m
 - Goal: use $O(\log_2 n)$ bits, $i_j \in \{1, \dots, n\}$

Randomized Approximate Counting

- Goal: compute estimate \tilde{n} of n where

$$P(|\tilde{n} - n| > \varepsilon \cdot n) < \delta$$

for some $0 < \varepsilon, \delta < 1$, that you can choose,
using $O(\log_2(\log_2 n))$ bits

- Chebyshev's Inequality:

$$P(|X - \mathbb{E}(X)| > \lambda) = P(|X - \mathbb{E}(X)|^2 > \lambda^2)$$

$$\leq \mathbb{E}(|X - \mathbb{E}(X)|^2) / \lambda^2 \quad \text{by Markov's Ineq.}$$

$$= \text{Var}(X) / \lambda^2 \quad \text{definition of Variance}$$

- X, Y independent $\Rightarrow \text{Var}(aX + bY) = a^2 \text{Var}(X) + b^2 \text{Var}(Y)$
- Ex: X_1, \dots, X_n independent, identically distributed (i.i.d.)

$$S = \frac{1}{n} \sum_{i=1}^n X_i \Rightarrow \mathbb{E}S = \mathbb{E}(X_i), \text{Var}(S) = \frac{1}{n} \text{Var}(X_i) \quad 3$$

First Try at Randomized Counting

Initialize: $c = 0$

Update: $c = c + 1$ with probability p

Query: return $\tilde{n} = c/p$

- Let $X_i = 1$ w.p. p , 0 w.p. $1-p$, i.i.d. $\Rightarrow c = \sum_{i=1}^n X_i$
- Thm: $\mathbb{E}(\tilde{n}) = \frac{1}{p} \mathbb{E}(c) = \frac{1}{p} \sum_{i=1}^n \mathbb{E}(X_i) = \frac{1}{p} \sum_{i=1}^n p = n$
- Thm: $\text{Var}(\tilde{n}) = \frac{1}{p^2} \text{Var}(c) = \frac{1}{p^2} \sum_{i=1}^n \text{Var}(X_i) = \frac{n}{p^2} (p - p^2) = \frac{n(1-p)}{p}$

• Do we save any bits?

#bits $= \lceil \log_2 c \rceil \approx \lceil \log_2 np \rceil \approx \log_2 n - \log_2 \frac{1}{p} \Rightarrow$ Save $\log_2 \frac{1}{p}$ bits

• Are we accurate? Not if we save many bits ($p \ll 1$)

$$\mathbb{P}(|n - \tilde{n}| > \varepsilon n) \leq \frac{\text{Var}(\tilde{n})}{(\varepsilon n)^2} = \frac{(1-p)}{\varepsilon^2 p n}$$

Morris's Algorithm:

(1/3)

Initialize: $X = 0$

Update: $X = X + 1$ with probability $\frac{1}{2^X}$

Query: return $\tilde{n} = 2^X - 1$

- Let $X_n = X$ after n updates
- Thm: $\mathbb{E}(\tilde{n}) = \mathbb{E}(2^{X_n} - 1) = n$, or $\mathbb{E}(2^{X_n}) = n + 1$

Proof: induction on n ; base case is $n = 0$

$$\mathbb{E}(2^{X_{n+1}}) = \sum_{j=0}^{\infty} P(X_n = j \text{ and we increment } X) \cdot 2^{j+1} \\ + P(X_n = j \text{ and we don't}) \cdot 2^j$$

$$= \sum_{j=0}^{\infty} P(X_n = j) \cdot \frac{1}{2^j} \cdot 2^{j+1} + P(X_n = j) \cdot \left(1 - \frac{1}{2^j}\right) 2^j$$

$$= \sum_{j=0}^{\infty} P(X_n = j) (2 + 2^j - 1) = \sum_{j=0}^{\infty} P(X_n = j) + \sum_{j=0}^{\infty} P(X_n = j) \cdot 2^j$$

$$= 1 + \mathbb{E}(2^{X_n}) = 1 + (n + 1) = n + 2$$

Morris's Algorithm:

(2/3)

Initialize: $X = 0$

Update: $X = X + 1$ with probability $\frac{1}{2^X}$

Query: return $\tilde{n} = 2^X - 1$

- Let $X_n = X$ after n updates
- Thm: $\mathbb{E}(\tilde{n}) = \mathbb{E}(2^{X_n} - 1) = n$
- Intuition: we are approximating $\log_2 n$ instead of $n \Rightarrow$ need $\log_2(\log_2 n)$ bits, exponentially fewer than $\log_2 n$ needed for n
- Could we use even fewer bits to approximate n to within a factor $1 \pm \epsilon$? No: need to distinguish $[1, (1+\epsilon)^2), [(1+\epsilon)^2, (1+\epsilon)^4) \dots [(1+\epsilon)^k, (1+\epsilon)^{k+2}), \dots [(1-\epsilon)n, (1+\epsilon)n]$
intervals $= \frac{\log_2 n (1+\epsilon)}{\log_2 (1+\epsilon)^2} = \Theta\left(\frac{\log_2 n}{\epsilon}\right) \Rightarrow$ need $\Omega(\log_2 \log_2 n)$ bits

Morris's Algorithm:

(3/3)

Initialize: $X = 0$

Update: $X = X + 1$ with probability $\frac{1}{2^X}$

Query: return $\tilde{n} = 2^X - 1$

- Let $X_n = X$ after n updates

- Thm: $\mathbb{E}(\tilde{n}) = \mathbb{E}(2^{X_n} - 1) = n$

- Thm: $\text{Var}(\tilde{n}) = \frac{1}{2}n^2 - \frac{1}{2}n - 1$

Proof: $\mathbb{E}(2^{2X_n}) = (3/2)n^2 + (3/2)n + 1$ by induction

$$\mathbb{E}(2^{2X_{n+1}}) = \sum_{j=0}^{\infty} P(X_n = j \text{ and we increment } X) \cdot 2^{2(j+1)}$$

$$+ P(X_n = j \text{ and we don't}) \cdot 2^{2j}$$

$$= \dots = (3/2)(n+1)^2 + (3/2)(n+1) + 1$$

$$\text{Var}(\tilde{n}) = \text{Var}(2^{X_n}) = \mathbb{E}(2^{2X_n}) - (\mathbb{E}(2^{X_n}))^2 = \dots = \frac{1}{2}n^2 - \frac{1}{2}n - 1$$

- Chebyshev: $P(|\tilde{n} - n| > \epsilon n) \leq \frac{1}{\epsilon^2 n^2} \frac{n^2}{2} = \frac{1}{2\epsilon^2}$. oops 7

Making Morris's Algorithm more accurate

- Run s "copies" of Morris, yielding $\tilde{n}_1, \dots, \tilde{n}_s$,
return average $\tilde{n} = \frac{1}{s} \sum_{i=1}^s \tilde{n}_i$
 - $\mathbb{E}(\tilde{n}) = \mathbb{E}(\tilde{n}_i) = n$, $\text{Var}(\tilde{n}) = \frac{1}{s} \text{Var}(\tilde{n}_i)$
- $P(|n - \tilde{n}| > \varepsilon n) \leq \frac{1}{s} \frac{1}{2\varepsilon^2} < \delta$ if $s > \frac{1}{2\varepsilon^2\delta} = \Theta\left(\frac{1}{\varepsilon^2\delta}\right)$
- How many bits do we need?
- Intuition: s copies of Morris $\Rightarrow O\left(\frac{1}{\varepsilon^2\delta} \cdot \log_2 \log_2 n\right)$ bits
- More carefully: with probability $1 - \delta$, need
 $O\left(\frac{1}{\varepsilon^2\delta} \cdot \log_2 \log_2 \frac{n}{\varepsilon\delta}\right)$ bits

Flajolet + Muller (FM) Alg. for Distinct Element Counting

- Given stream i_1, i_2, \dots, i_m , each $i_j \in \{1, \dots, n\}$
count $t = \#$ distinct elements in stream

• Ex if stream is $1, 2, 7, 2, 3, 7$, $t = 4$

- Straightforward solutions:

1) Keep array of n bits y_j , initially all $y_j = 0$
set $y_{i_k} = 1$ when i_k appears


2) Store whole stream in memory $\Rightarrow O(m \cdot \log_2 n)$ bits

- Goal: use $o(n)$ bits to compute \tilde{t}

where $P(|t - \tilde{t}| > \varepsilon t) < \delta$

Idealized FM Alg.

- Goal: count $t = \#$ distinct elements in $i_1, \dots, i_m, i_j \in \{1..n\}$
- Pick random function $h: \{1, \dots, n\} \rightarrow [0, 1]$
 - Each $h(i)$ is i.i.d. random real number uniformly distributed in $[0, 1]$
- Initialize: $X = 1$
Update: $X = \min(X, h(i))$
Query: return $\tilde{t} = \frac{1}{X} - 1$
- Intuition: $X = \min$ of t distinct uniform random i.i.d. numbers in $[0, 1]$, we expect $X \sim \frac{1}{t+1}$, so $t \sim \frac{1}{X} - 1$



Idealized FM Alg.

- Goal: count $t = \#$ distinct elements in $i_1, \dots, i_m, i_j \in \{1..n\}$
- $X = \min$ of t uniform random i.i.d numbers in $[0,1]$
- Thm: $\mathbb{E}(X) = \frac{1}{t+1}$

Proof: Analogous to discrete case where

$$\mathbb{E}(X) = \sum_{i=1}^{\infty} i p(i) = \sum_{i=1}^{\infty} \sum_{j=i}^{\infty} p(j) = \sum_{i=1}^{\infty} P(X \geq i)$$

$$\mathbb{E}(X) = \int_0^{\infty} P(X > \lambda) d\lambda = \int_0^1 P(X > \lambda) d\lambda$$

$$= \int_0^1 P(\text{all } t \text{ random \#s are } > \lambda) d\lambda$$

$$= \int_0^1 \prod_{i=1}^t P(\text{one random \# is } > \lambda) d\lambda$$

$$= \int_0^1 \prod_{i=1}^t (1 - \lambda) d\lambda = \int_0^1 (1 - \lambda)^t d\lambda = \frac{1}{t+1}$$

Making Idealized FM More Accurate

- Same as Morris: run s copies, average results

- Thm: $\text{Var}(X) = \frac{t}{(t+1)^2(t+2)}$

Proof: Follows from

$$\begin{aligned}\mathbb{E}(X^2) &= \int_0^1 P(X^2 > \lambda) d\lambda = \int_0^1 P(X > \sqrt{\lambda}) d\lambda \\ &= \int_0^1 (1 - \sqrt{\lambda})^t d\lambda \quad \text{substitute } u = 1 - \sqrt{\lambda} \\ &= 2 \int_0^1 u^t (1-u) du = \frac{2}{(t+1)(t+2)}\end{aligned}$$

- Thm: run $s = \lceil \frac{1}{\epsilon^2 \delta} \rceil$ independent copies FM_1, \dots, FM_s with outputs X_1, \dots, X_s , $Z = \frac{1}{s} \sum_{i=1}^s X_i$, output $\tilde{t} = 1/Z - 1$
Then $P(|\tilde{t} - t| > O(\epsilon)t) \leq \delta$

Proof: Apply Chebyshev's Ineq. to Z

Making FM Practical

- Can't generate uniform random real numbers in practice, need an approximation
- Generate random integers $\in \{0, 1, \dots, B\}$, divide final minimum by B
- Use "pseudorandom number generators" to approximate actual random numbers
 - Uses $O(\log n + \log B)$ bits
- Recent version: Hyperloglog