

Lecture #24

CS170

Spring 2021



Lower Bounds

- Complexity of a problem P is a function $T(P)$ that measures its cost (time/memory/...) as a function $f(n)$ of its input size n
- An algorithm for P gives an upper bound on $f(n)$
 - Eg $P = \text{sorting}$, $T(P) = \# \text{ comparisons}$
 - Insertion Sort $\Rightarrow f(n) \leq O(n^2)$
 - Merge Sort $\Rightarrow f(n) \leq O(n \log n)$
- A lower bound for P is a proof that $f(n) = \Omega(g(n))$
 - Holds for any algorithm (in a class)
 - Eg: $f(n) = \Omega(n \log n)$ for $\# \text{ comparisons}$ in sorting
 - $2^{\# \text{ comparisons}} \geq n! \geq (n/2)^{(n/2)}$ (not bucket sort)

Lower Bounds

- Recall NP-complete problems: 3SAT, ILP, ...
 - widely believed lower bound: $\Omega(n^{w(1)})$
i.e. bigger than any polynomial
 - Best known lower bounds: $\Omega(n)$
 - Proving lower bounds is hard
- CS 172 preview
 - "Time Hierarchy Thm"
 - Given code implementing Boolean function C , that takes a binary string X of length $|X|=n$ as input, will C return true in $\leq n^3$ steps?
 - Naive algorithm: run it for n^3 steps and see
 - Thm: Any correct algorithm takes $\Omega(n^{3-\epsilon})$ steps for some tiny ϵ

Examples of Lower Bounds for specific classes of algorithms

1) Circuit Complexity:

- What size circuit (#wires or depth) is needed to solve a problem of input size n ?

2) Cell Probe Model

- How many reads/writes to memory are needed to solve a problem of size n ?

3) Branching Program

- How do time and memory needs trade off?


4) Communication Complexity

- If Proc 1 owns X , and Proc 2 owns Y , how many bits do they need to exchange to compute $f(X, Y)$?

Circuit Complexity (1/3)

- Problem: Given $f: \{0,1\}^n \rightarrow \{0,1\}$, how big a circuit do you need to evaluate f ?
 - Circuit = DAG of and, or, not gates
 - Size could be # wires, depth
- Ex $f: \{0,1\}^{10} \rightarrow \{0,1\}^{10}$ could multiply $x \cdot y$ where x = first 5 bits of input, y = last 5 bits
- What is known?
 - # circuits = $2^{\Theta(w \log w)}$, w = # wires
 - # functions on n input bits = 2^{2^n}
 - $2^{C w \log w} \geq 2^{2^n} \Rightarrow C \cdot w \cdot \log w \geq 2^n$
 - Even if $w \sim 1.9^n$ most functions need more wires
 - Most $f: \{0,1\}^n \rightarrow \{0,1\}$ need exponentially many wires

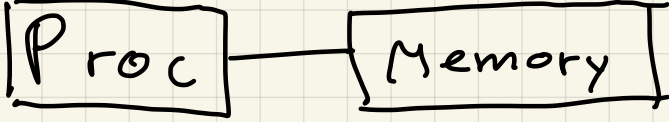
Circuit Complexity (2/3)

- Problem: Given $f: \{0,1\}^n \rightarrow \{0,1\}$, how big a circuit do you need to evaluate f ?
 - Circuit = DAG of and, or, not gates
 - Size could be # wires, depth
- Most $f: \{0,1\}^n \rightarrow \{0,1\}$ need exponentially many wires
 - Do we know any?
 - Best result so far (2016): $\# \text{wires} \geq (3 + \frac{1}{86}) \cdot n$
- What about depth?
 - $f = \text{Parity} = \text{XOR}$ of all n input bits
 - Compute f using binary tree of depth $\log_2 n$
 - What if we restrict depth to a constant k ?
and allow unbounded fan in: 
 - Thm: need $\# \text{wires} = e^{\Omega(n)}$

Circuit Complexity (3/3)


- Problem: Given $f: \{0,1\}^n \rightarrow \{0,1\}$, how big a circuit do you need to evaluate f ?
 - Circuit = DAG of and, or, not gates
 - Size could be # wires, depth
- Connection to NP-Completeness
 - Def: $P/poly$ = set of all problems that can be solved with circuit of $\#wires \leq poly(n)$
 - $P/poly$ is analogue of P for circuit complexity
 - Known: $P \subseteq P/poly$
 - Unknown: is $NP \subseteq P/poly$? If not, $P \neq NP$

Cell Probe Model

- Algorithm is allowed to perform the following ops:
 - A diagram showing a box labeled 'Proc' connected by a horizontal line to a box labeled 'Memory'.
 - Processor can read a word (w bits) from memory location i , or write a word
 - How many reads/writes to memory are needed to solve a problem of size n ?
- Used to find lower bounds on cost of using data structures

Branching programs

- DAG to compute $f: \{0,1\}^n \rightarrow Y$
 - one source node
 - one sink node per element of Y
 - each non-sink nodes has 2 outgoing edges labelled $x_i=0$ or $x_i=1$, where $x=\text{input}$
- Adding 3 bits: $Y = \{0,1,2,3\}$



Width = 4 = 2^{#bits needed}
#Layers = 3 = #steps
to compute
answer
≈ runtime

- How do Width(#bits) and #Layers(runtime) trade off?

Communication Complexity (1/8)

- Alice and Bob both want $f:(X,Y) \rightarrow \{0,1\}$
 - Alice only knows X , Bob only knows Y
 - They exchange messages, last one to receive a message announces $f(X,Y)$
 - Goal: minimize #bits Alice and Bob exchange
- Different kinds of algorithms allowed:
 - Deterministic
 - Public coin randomness: Alice and Bob have same random bits
 - Private coin randomness: Alice and Bob have their own random bits
 - One way communication (Alice sends one message to Bob) vs. 2-way

Communication Complexity (2/8)

- Alice knows X , Bob knows Y , want $f(X, Y) \in \{0, 1\}$ while minimizing #bits exchanged
- Def: $D(f)$ = minimum #bits with deterministic alg
- Def: $R_{\text{pub}}(f)$ = minimum #bits with **public** coin randomness
- Def: $R_{\text{priv}}(f)$ = minimum #bits with **private** coin randomness
- Thm: $D(f) \geq R_{\text{priv}}(f) \geq R_{\text{pub}}(f)$
Proof: 1st ineq: deterministic special case: ignore random bits
2nd ineq: Alice uses 1st half of random bit, Bob 2nd half
- Def: $EQ(X, Y) = 1$ if $X=Y$ else 0
- Thm: $D(EQ) = \Theta(n)$, $R_{\text{priv}}(EQ) = \Theta(\log n)$, $R_{\text{pub}}(EQ) = \Theta(1)$
- Simplify: consider one-way communication cost, $D^{\rightarrow}(f)$,
Alice sends one message to Bob

Communication Complexity (3/8)

- Alice knows X , Bob knows Y , want $f(X, Y) \in \{0, 1\}$
- Def: $D(f)$ = minimum #bits with deterministic alg where Alice sends one message to Bob
- Claim: $D(EQ) \geq n = |X|$
proof: if Alice sends $g(x) \in \{0, 1\}^m$ to Bob with $m < n$,
 $\exists X_1 \neq X_2$ but $g(X_1) = g(X_2)$ so Bob can't tell the difference between X_1 and X_2
- Def: DE = counting #distinct elements
- Claim: Any exact deterministic alg A for DE requires $\Omega(n)$ bits of memory (FM was random)
proof: Show that if A solves DE with s bits, we can use A to solve EQ with $s + \log n$ bits, so
 $s + \log n \geq D(EQ) \geq n \Rightarrow s \geq n - \log n = \Omega(n)$

Communication Complexity (4/8)

- Alice knows X , Bob knows Y , want $f(X, Y) \in \{0, 1\}$
 - Def: $D^{\rightarrow}(f)$ = minimum #bits with deterministic alg where Alice sends one message to Bob
 - Claim: $D^{\rightarrow}(EQ) \geq n = |X|$
 - Def: DE = counting #distinct elements
 - Claim: Any exact deterministic alg A for DE requires $\Omega(n)$ bits of memory (FM was random)
- proof: Show that if A solves DE with s bits, we can use A to solve EQ with $s + \log n$ bits, so
- $$s + \log n \geq D^{\rightarrow}(EQ) \geq n \Rightarrow s \geq n - \log n = \Omega(n)$$

How to use A to solve DE : For each $X_i = 1$, Alice feeds i into A , sends $(\text{mem}(A), \#1s \text{ in } X) = s + \log n$ bits to Bob. Bob checks if $\#1s \text{ in } X = \#1s \text{ in } Y$. If yes, Bob asks A for $\#DE$, then feeds i for each $Y_i = 1$ into A and asks if $\#DE$ increases. If not, $X = Y$, else $X \neq Y$.

Communication Complexity (5/8)

- Claim: Any exact deterministic alg A for DE requires $\Omega(n)$ bits of memory
- Claim: Any approximate deterministic alg for DE ($|t - \tilde{t}| \leq 0.01t$) also requires $\Omega(n)$ bits of memory

Proof: consider EQ where $X, Y \in B \subseteq \{0, 1\}^n$

Same argument as before $\Rightarrow D^{\rightarrow}(EQ^B) \geq \log_2 |B|$

Claim: (no proof!) $\exists B$ such that

$|B| \geq 2^{cn}$, all $X \in B$ have same $\#1s = r$,

and if $X, Y \in B$, $X \neq Y$, then $\#1s$ X and Y share $\leq \frac{r}{10}$

(Think of $X, Y \subseteq \{1, \dots, n\}$, so $|X| = |Y| = r$, $|X \cap Y| \leq \frac{r}{10}$)

Show that if A_{ap} solves DE with s bits, 1% error,

we can use A_{ap} to solve EQ^B with s bits

$\Rightarrow s \geq D^{\rightarrow}(EQ^B) \geq \log_2 |B| \geq cn = \Omega(n)$

Communication Complexity (6/8)

- Claim: Any approximate deterministic alg for DE ($|t - \tilde{t}| \leq 0.01t$) also requires $\Omega(n)$ bits of memory

Proof: EQ where $X, Y \in B \subseteq \{0,1\}^n$, $D^{\rightarrow}(EQ^B) \geq \log_2 |B|$

Claim: (no proof!) $\exists B$ such that

$|B| \geq 2^{cn}$, all $X \in B$ have same $\#1s = r$,

and if $X, Y \in B$, $X \neq Y$, then $\#1s$ X and Y share $\leq \frac{r}{10}$

Show that if A_{ap} solves DE with s bits, 1% error,

we can use A_{ap} to solve EQ^B with s bits

$$\Rightarrow s \geq D^{\rightarrow}(EQ^B) \geq \log_2 |B| \geq cn = \Omega(n)$$

How to use A_{ap} to solve EQ^B (similar to using A for EQ):

For each $X_i = 1$, Alice feeds i into A_{ap} , sends $\text{mem}(A_{ap})$ to Bob.

Bob knows $\#DE = r$ (property of B). For each $Y_i = 1$, Bob feeds i into A_{ap} , asks for $\#DE$. 2 cases:

$X = Y \Rightarrow A_{ap}$ reports $\#DE \in [.99r, 1.01r]$



$X \neq Y \Rightarrow A_{ap}$ reports $\#DE \in [.99 \cdot 1.9 \cdot r, 1.01 \cdot 2 \cdot r]$] different!

Communication Complexity (7/8)

- Summary of counting distinct elements (DE)
 - No exact, deterministic alg with $o(n)$ memory
 - No approx, deterministic alg with $o(n)$ memory
 - No exact, randomized alg with $o(n)$ memory
 - \Rightarrow need approx, randomized alg to use $o(n)$ memory (FM)
- How to show B exists? choose randomly, show it has right property with probability > 0

Communication Complexity (8/8)

- Goal: minimize communication between main memory and cache, or between processors connected over a network



- Thm (Hong, Kung, 81) Any execution of $\Theta(n^3)$ matrix multiply moves $\Omega(n^3/\sqrt{M})$ words between Cache and main memory
- Attained by "loop tiling", widely implemented
- Extends to rest of linear algebra, any code that looks like nested loops accessing arrays